

```
from google.colab import drive
drive.mount('/content/gdrive')
```

Mounted at /content/gdrive

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
from google.colab import files
data_to_load=files.upload()
```

<IPython.core.display.HTML object>

Saving loan\_data\_set.csv to loan\_data\_set.csv

```
df=pd.read_csv('loan_data_set.csv')
df
```

	Loan_ID	Gender	Married	...	Credit_History	Property_Area
Loan_Status						
0	LP001002	Male	No	...	1.0	Urban
Y						
1	LP001003	Male	Yes	...	1.0	Rural
N						
2	LP001005	Male	Yes	...	1.0	Urban
Y						
3	LP001006	Male	Yes	...	1.0	Urban
Y						
4	LP001008	Male	No	...	1.0	Urban
Y						
...	...	...	...	...	...	...
...						
609	LP002978	Female	No	...	1.0	Rural
Y						
610	LP002979	Male	Yes	...	1.0	Rural
Y						
611	LP002983	Male	Yes	...	1.0	Urban
Y						
612	LP002984	Male	Yes	...	1.0	Urban
Y						
613	LP002990	Female	No	...	0.0	Semiurban
N						

[614 rows x 13 columns]

```
df.isnull().sum()
```

Loan_ID	0
Gender	13
Married	3


```

Dependents      15
Education        0
Self_Employed   32
ApplicantIncome  0
CoapplicantIncome 0
LoanAmount      22
Loan_Amount_Term 14
Credit_History  50
Property_Area    0
Loan_Status      0
dtype: int64

```

```
sns.heatmap(df.isnull())
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7feb6710f2d0>
```



```
print(df.info())
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
 #   Column              Non-Null Count  Dtype
---  -

```

0	Loan_ID	614	non-null	object
1	Gender	601	non-null	object
2	Married	611	non-null	object
3	Dependents	599	non-null	object
4	Education	614	non-null	object
5	Self_Employed	582	non-null	object
6	ApplicantIncome	614	non-null	int64
7	CoapplicantIncome	614	non-null	float64
8	LoanAmount	592	non-null	float64
9	Loan_Amount_Term	600	non-null	float64
10	Credit_History	564	non-null	float64
11	Property_Area	614	non-null	object
12	Loan_Status	614	non-null	object

dtypes: float64(4), int64(1), object(8)

memory usage: 62.5+ KB

None

df.nunique()

Loan_ID	614
Gender	2
Married	2
Dependents	4
Education	2
Self_Employed	2
ApplicantIncome	505
CoapplicantIncome	287
LoanAmount	203
Loan_Amount_Term	10
Credit_History	2
Property_Area	3
Loan_Status	2

dtype: int64

data\_numerical =

df[['ApplicantIncome', 'CoapplicantIncome', 'LoanAmount', 'Loan\_Amount\_Term', 'Credit\_History']]

data\_numerical.columns

Index(['ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',  
'Loan\_Amount\_Term', 'Credit\_History'],  
dtype='object')

data\_non\_numerical=df[['Loan\_ID', 'Gender', 'Married', 'Dependents', 'Education', 'Self\_Employed', 'Property\_Area', 'Loan\_Status']]

data\_non\_numerical.columns

```
Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
      'Self_Employed', 'Property_Area', 'Loan_Status'],
      dtype='object')
```

```
data_numerical.head()
```

```

ApplicantIncome  CoapplicantIncome  ...  Loan_Amount_Term
Credit_History
0                5849                0.0  ...              360.0
1.0
1                4583              1508.0  ...              360.0
1.0
2                3000                0.0  ...              360.0
1.0
3                2583              2358.0  ...              360.0
1.0
4                6000                0.0  ...              360.0
1.0
```

```
[5 rows x 5 columns]
```

```
data_non_numerical.head()
```

```

Loan_ID Gender Married  ... Self_Employed Property_Area
Loan_Status
0 LP001002  Male      No  ...              No      Urban
Y
1 LP001003  Male     Yes  ...              No      Rural
N
2 LP001005  Male     Yes  ...             Yes      Urban
Y
3 LP001006  Male     Yes  ...              No      Urban
Y
4 LP001008  Male      No  ...              No      Urban
Y
```

```
[5 rows x 8 columns]
```

```
data_non_numerical.groupby('Gender')['Loan_ID'].count()
```

```

Gender
Female    112
Male      489
Name: Loan_ID, dtype: int64
```

```
x=data_non_numerical.columns
```

```
for i in x[1:]:
```

```
    m=data_non_numerical.groupby(i)['Loan_ID'].count().idxmax()
```

```
    data_non_numerical[i].fillna(m,inplace=True)
```

```

/usr/local/lib/python3.7/dist-packages/pandas/core/series.py:4536:
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

downcast=downcast,

```
data_non_numerical.isnull().sum()
```

```
Loan_ID      0
Gender       0
Married      0
Dependents   0
Education    0
Self_Employed 0
Property_Area 0
Loan_Status  0
dtype: int64
```

```
data_non_numerical
```

	Loan_ID	Gender	Married	...	Self_Employed	Property_Area
Loan_Status						
0	LP001002	Male	No	...	No	Urban
Y						
1	LP001003	Male	Yes	...	No	Rural
N						
2	LP001005	Male	Yes	...	Yes	Urban
Y						
3	LP001006	Male	Yes	...	No	Urban
Y						
4	LP001008	Male	No	...	No	Urban
Y						
..	...	...	...	...	...	...
...						
609	LP002978	Female	No	...	No	Rural
Y						
610	LP002979	Male	Yes	...	No	Rural
Y						
611	LP002983	Male	Yes	...	No	Urban
Y						
612	LP002984	Male	Yes	...	No	Urban
Y						
613	LP002990	Female	No	...	Yes	Semiurban
N						

```
[614 rows x 8 columns]
```

```
data_numerical.columns
```

```
Index(['ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',  
      'Loan_Amount_Term', 'Credit_History'],  
      dtype='object')
```

```
x=data_numerical.columns
```

```
for i in x:
```

```
    m=data_numerical[i].mean()
```

```
    data_numerical[i].fillna(m, inplace=True)
```

```
/usr/local/lib/python3.7/dist-packages/pandas/core/series.py:4536:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
    downcast=downcast,
```

```
data_numerical.isnull().sum()
```

```
ApplicantIncome      0
```

```
CoapplicantIncome    0
```

```
LoanAmount           0
```

```
Loan_Amount_Term     0
```

```
Credit_History       0
```

```
dtype: int64
```

```
data_numerical
```

	ApplicantIncome	CoapplicantIncome	...	Loan_Amount_Term
Credit_History				
0	5849	0.0	...	360.0
1.0				
1	4583	1508.0	...	360.0
1.0				
2	3000	0.0	...	360.0
1.0				
3	2583	2358.0	...	360.0
1.0				
4	6000	0.0	...	360.0
1.0				
...	...	...	...	...
...				
609	2900	0.0	...	360.0
1.0				
610	4106	0.0	...	180.0
1.0				
611	8072	240.0	...	360.0
1.0				
612	7583	0.0	...	360.0
1.0				

```
613          4583          0.0 ...          360.0
0.0
```

```
[614 rows x 5 columns]
```

```
data_non_numerical.isnull().sum()
```

```
Loan_ID      0
Gender        0
Married       0
Dependents    0
Education     0
Self_Employed 0
Property_Area 0
Loan_Status   0
dtype: int64
```


```
data_non_numerical
```

	Loan_ID	Gender	Married	...	Self_Employed	Property_Area
Loan_Status						
0	LP001002	Male	No	...	No	Urban
Y						
1	LP001003	Male	Yes	...	No	Rural
N						
2	LP001005	Male	Yes	...	Yes	Urban
Y						
3	LP001006	Male	Yes	...	No	Urban
Y						
4	LP001008	Male	No	...	No	Urban
Y						
..	...	...	...	...	...	...
...						
609	LP002978	Female	No	...	No	Rural
Y						
610	LP002979	Male	Yes	...	No	Rural
Y						
611	LP002983	Male	Yes	...	No	Urban
Y						
612	LP002984	Male	Yes	...	No	Urban
Y						
613	LP002990	Female	No	...	Yes	Semiurban
N						

```
[614 rows x 8 columns]
```

```
sns.heatmap(data_numerical.isnull())
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7feb5e27c310>
```



```
data_numerical.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 614 entries, 0 to 613
```

```
Data columns (total 5 columns):
```

#	Column	Non-Null Count	Dtype
0	ApplicantIncome	614 non-null	int64
1	CoapplicantIncome	614 non-null	float64
2	LoanAmount	614 non-null	float64
3	Loan_Amount_Term	614 non-null	float64
4	Credit_History	614 non-null	float64

```
dtypes: float64(4), int64(1)
```

```
memory usage: 24.1 KB
```

```
data_non_numerical.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 614 entries, 0 to 613
```

```
Data columns (total 8 columns):
```

#	Column	Non-Null Count	Dtype
0	Loan_ID	614 non-null	object



```

1  Gender          614 non-null    object
2  Married          614 non-null    object
3  Dependents       614 non-null    object
4  Education        614 non-null    object
5  Self_Employed    614 non-null    object
6  Property_Area    614 non-null    object
7  Loan_Status      614 non-null    object

```

```
dtypes: object(8)
```

```
memory usage: 38.5+ KB
```

```
data_non_numerical.columns
```

```
Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
      'Self_Employed', 'Property_Area', 'Loan_Status'],
      dtype='object')
```

```
dummies=pd.get_dummies(data_non_numerical,columns=['Gender','Married',
'Dependents',
      'Education', 'Self_Employed', 'Property_Area',
'Loan_Status'])
```

```
dummies
```

	Loan_ID	Gender_Female	...	Loan_Status_N	Loan_Status_Y
0	LP001002	0	...	0	1
1	LP001003	0	...	1	0
2	LP001005	0	...	0	1
3	LP001006	0	...	0	1
4	LP001008	0	...	0	1
..	...	...	...	...	...
609	LP002978	1	...	0	1
610	LP002979	0	...	0	1
611	LP002983	0	...	0	1
612	LP002984	0	...	0	1
613	LP002990	1	...	1	0

```
[614 rows x 18 columns]
```

```
from sklearn.preprocessing import StandardScaler
```

```
sc=StandardScaler()
```

```
data_numerical.iloc[:,1:]=sc.fit_transform(data_numerical.iloc[:,1:])
```

```
/usr/local/lib/python3.7/dist-packages/pandas/core/indexing.py:1734:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation:
```

```
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#
returning-a-view-versus-a-copy
```

```
    isetter(loc, value[:, i].tolist())
```

```
#data_numerical.iloc[2:8,2:5]
```

	LoanAmount	Loan_Amount_Term	Credit_History
2	-0.957641	0.279851	0.451640
3	-0.314547	0.279851	0.451640
4	-0.064454	0.279851	0.451640
5	1.436099	0.279851	0.451640
6	-0.612275	0.279851	0.451640
7	0.138001	0.279851	-2.410441

```
df_loan=pd.concat([data_numerical.iloc[:,1:],dummies.drop(['Loan_ID'],  
axis=1)],axis=1)
```

```
df_loan
```

	CoapplicantIncome	LoanAmount	...	Loan_Status_N	Loan_Status_Y
0	-0.554487	0.000000	...	0	1
1	-0.038732	-0.219273	...	1	0
2	-0.554487	-0.957641	...	0	1
3	0.251980	-0.314547	...	0	1
4	-0.554487	-0.064454	...	0	1
...	...	...	...	...	...
609	-0.554487	-0.898095	...	0	1
610	-0.554487	-1.267279	...	0	1
611	-0.472404	1.269371	...	0	1
612	-0.554487	0.483367	...	0	1
613	-0.554487	-0.159728	...	1	0

```
[614 rows x 21 columns]
```

```
df_loan.isnull().sum()
```

CoapplicantIncome	0
LoanAmount	0
Loan_Amount_Term	0
Credit_History	0
Gender_Female	0
Gender_Male	0
Married_No	0
Married_Yes	0
Dependents_0	0
Dependents_1	0
Dependents_2	0
Dependents_3+	0
Education_Graduate	0
Education_Not Graduate	0
Self_Employed_No	0
Self_Employed_Yes	0
Property_Area_Rural	0

```
Property_Area_Semiurban    0
Property_Area_Urban        0
Loan_Status_N              0
Loan_Status_Y              0
dtype: int64
```

```
from sklearn.model_selection import train_test_split
x=df_loan.drop(["Loan_Status_N","Loan_Status_Y"],axis=1)
y=df_loan["Loan_Status_Y"]
X_train,X_test,y_train,y_test
=train_test_split(x,y,test_size=0.2,random_state=30)
```

X\_test

	CoapplicantIncome	...	Property_Area_Urban
193	-0.554487	...	0
334	0.072423	...	1
343	-0.554487	...	0
24	0.445901	...	0
47	-0.554487	...	1
..	...	...	...
453	0.288917	...	0
31	-0.554487	...	1
293	-0.554487	...	1
410	-0.554487	...	0
108	0.676760	...	1

[123 rows x 19 columns]

y\_test

193	1
334	1
343	1
24	0
47	1
..	..
453	1
31	0
293	0
410	0
108	0

Name: Loan\_Status\_Y, Length: 123, dtype: uint8

```
from sklearn.svm import SVC
svclassifier=SVC(kernel='linear')
svclassifier.fit(X_train,y_train)
```

```
SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None,
coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='scale',
kernel='linear',
```

```
max_iter=-1, probability=False, random_state=None, shrinking=True,  
tol=0.001, verbose=False)
```

```
pred_1=svclassifier.predict(X_test)
```

```
from sklearn.metrics import classification_report, confusion_matrix,  
accuracy_score  
print(confusion_matrix(y_test,pred_1))  
print(classification_report(y_test,pred_1))  
print(accuracy_score(y_test,pred_1))
```

```
[[13 24]  
 [ 2 84]]
```

	precision	recall	f1-score	support
0	0.87	0.35	0.50	37
1	0.78	0.98	0.87	86
accuracy			0.79	123
macro avg	0.82	0.66	0.68	123
weighted avg	0.80	0.79	0.76	123

```
0.7886178861788617
```

```
from sklearn.neighbors import KNeighborsClassifier  
knn = KNeighborsClassifier()  
knn.fit(X_train, y_train)
```

```
KNeighborsClassifier(algorithm='auto', leaf_size=30,  
metric='minkowski',  
metric_params=None, n_jobs=None, n_neighbors=5,  
p=2,  
weights='uniform')
```

```
pred_2=knn.predict(X_test)
```

```
from sklearn.metrics import classification_report, confusion_matrix,  
accuracy_score  
print(confusion_matrix(y_test,pred_2))  
print(classification_report(y_test,pred_2))  
print(accuracy_score(y_test,pred_2))
```

```
[[14 23]  
 [ 7 79]]
```

	precision	recall	f1-score	support
0	0.67	0.38	0.48	37
1	0.77	0.92	0.84	86
accuracy			0.76	123
macro avg	0.72	0.65	0.66	123

weighted avg            0.74            0.76            0.73            123

0.7560975609756098

```
from sklearn.naive_bayes import GaussianNB
clf = GaussianNB()
clf.fit(X_train, y_train)
pred_3=clf.predict(X_test)
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score
print(confusion_matrix(y_test,pred_3))
print(classification_report(y_test,pred_3))
print(accuracy_score(y_test,pred_3))
```

```
[[14 23]
 [ 5 81]]
```

	precision	recall	f1-score	support
0	0.74	0.38	0.50	37
1	0.78	0.94	0.85	86
accuracy			0.77	123
macro avg	0.76	0.66	0.68	123
weighted avg	0.77	0.77	0.75	123

0.7723577235772358

```
from sklearn.tree import DecisionTreeClassifier
dtree = DecisionTreeClassifier(random_state=0)
tree = dtree.fit(X_train, y_train)
pred_4=tree.predict(X_test)

from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score
print(confusion_matrix(y_test,pred_4))
print(classification_report(y_test,pred_4))
print(accuracy_score(y_test,pred_4))
```

```
[[19 18]
 [20 66]]
```

	precision	recall	f1-score	support
0	0.49	0.51	0.50	37
1	0.79	0.77	0.78	86
accuracy			0.69	123
macro avg	0.64	0.64	0.64	123
weighted avg	0.70	0.69	0.69	123

0.6910569105691057

```
from sklearn.ensemble import VotingClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import AdaBoostClassifier,
RandomForestClassifier, ExtraTreesClassifier
from sklearn.metrics import accuracy_score, f1_score, log_loss

voting_clf=VotingClassifier(estimators=[('SVC', svcclassifier),
('DTree',dtree),('NB',clf),('KNN',knn)])
voting_clf.fit(X_train, y_train)
preds=voting_clf.predict(X_test)
acc=accuracy_score(y_test, preds)
l_loss=log_loss(y_test, preds)
f1=f1_score(y_test, preds)

print("Accuracy is: "+str(acc))
print("Log Loss is: "+str(l_loss))
print("F1 Score is: "+str(f1))

Accuracy is: 0.6910569105691057
Log Loss is: 10.670626798072515
F1 Score is: 0.7738095238095238
```